



# Building Modern and Real-Time Applications

Using Azure Cache for Redis Enterprise and Azure Spring Cloud



Sean Noyes - Senior Cloud Partner Solution Architect @Redis

[sean.noyes@redis.com](mailto:sean.noyes@redis.com)   
 @seanbot2000



# Agenda

## Today's topics

- A short history of the NoSQL movement
- Modern application architecture and Redis Use Cases
- Azure Cache for Redis Enterprise
- An example of modern data-centric architecture
- Let's build an environment and deploy something!
- Q&A

# A short history of the NoSQL movement

# A short history of the NoSQL movement

What's a NoSQL anyway?

- NoSQL today stands for ‘Not only SQL’
- Whereby SQL is used as a synonym for Relational Database Management Systems (RDBMS)
- First coined in 1998 by Carlo Strozzi, who built an open source relational database without a SQL interface
- It’s a movement, not a standard (unlike SQL, which *is* a standard)

# A short history of the NoSQL movement

## Why the need for non-relational databases?

### Not-Only SQL

- Optimized for specific use cases
- Flexible schemas
  - Can change on-demand, logic moved into app tier
- Scalability & performance
  - Scale out instead of scale up
- Availability
  - Designed from the ground up to be distributed and always available.....but at the sacrifice of.....
- Eventual consistency
  - With varying levels of strength, usually not ACID compliant
- Cloud Native from the ground up
  - Born in the modern age
- Simplified administration
  - Ops/DevOps instead of DBAs

### SQL (RDBMS)

- Aim to be general purpose systems
  - Often good enough at most cases, but not very well optimised for specific ones
- Fixed schemas
  - Changes often mean migrations
- Typically hard to scale (but might be good enough)
  - Favoring a scale-up over a scale-out approach
  - Which limits scalability
  - And availability at scale, esp. geographically
- But they are consistent (ACID compliance)
- Complex administration
  - Need for DBAs
- Though there are a lot of new cool kids on the block!
  - Amazon RDS, Aurora, Azure SQL
  - Yugabyte

# A short history of the NoSQL movement

## Specific use cases for NoSQL

- Key-value store
  - Memcached
- Data structure store
  - Redis
- Wide-table store
  - Cassandra
- Document database
  - MongoDB, Cosmos DB
- Graph database
  - Neo4J
- Time series database
  - InfluxDB



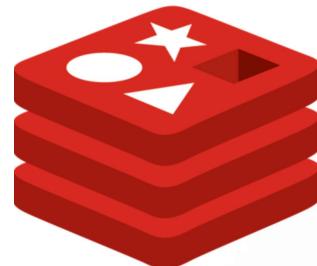
# Redis in a nutshell

*NoSQL Database, Cache,  
Message Broker*

*Key-value and  
data structure  
store*

*In-Memory  
Architecture*

*Written in C*



# redis

Remote Dictionary Server  
Open-Source (BSD License)

*Multiple  
Persistence  
Options*

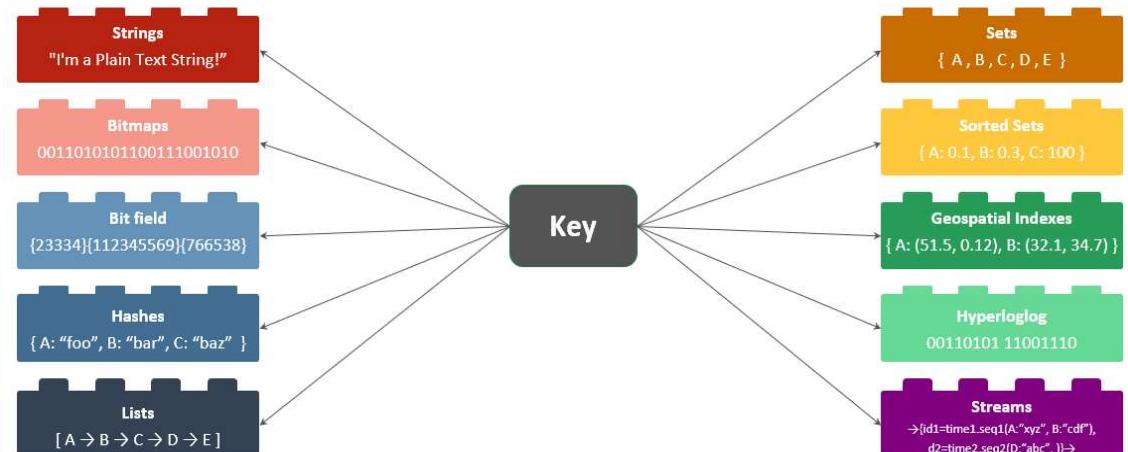
*ACID Compliant*

*Single Threaded*

# The core principles of Redis

## Three pillars

- Performance
  - All data in-memory
  - Single-threaded, lock-free
  - Most commands O(1) complexity
  - Efficient communication protocols
- Simplicity
  - Key -> Data structure
  - Data structure = Strings, (Sorted) Sets, Lists, Geospatial indexes, Hashes, Bit fields/maps, Pub/Sub, HyperLogLog and Streams
  - Extensible via modules, e.g. time series, graphs, search, etc.
- Use-case driven design
  - Data structures map to specific use cases, e.g. caching, user sessions, location services, recommendations, leaderboards, etc.



# Modern, real-time data architectures (with Azure Cache for Redis)

## USE CASE 1

# Distributed cache

## Use case examples

Handling spikes in traffic during major events

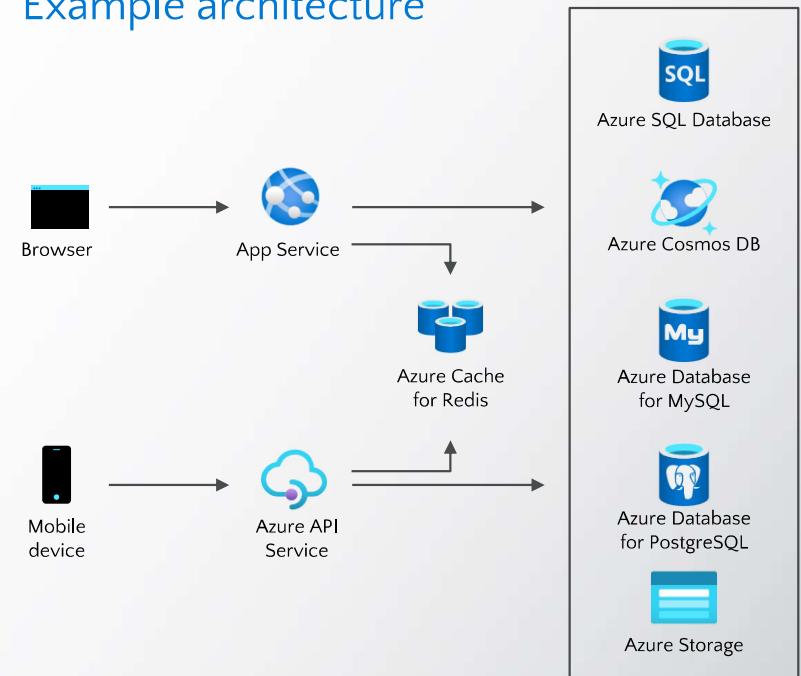
Serving commonly accessed data to users, like inventory or pricing data

Reducing compute load on a database

Placing content geographically closer to a consumer

Output caching

## Example architecture



## USE CASE 2

# Session store

### Use case examples

eCommerce shopping carts

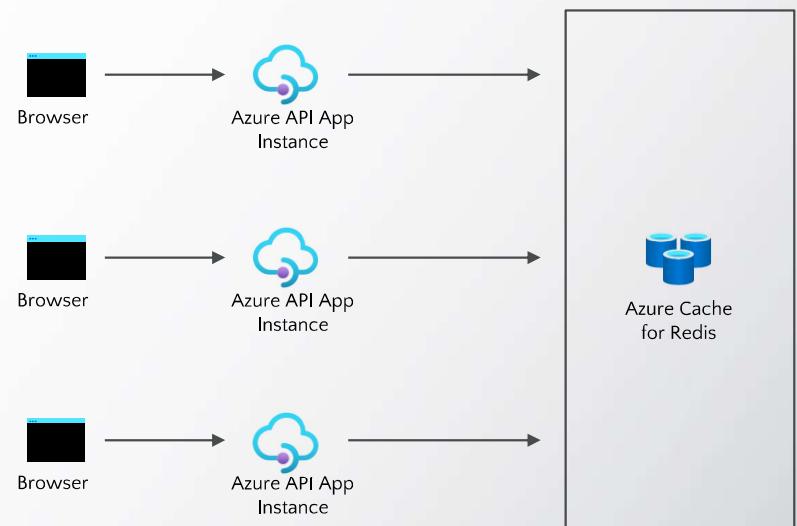
Storing user cookies

User login and session state

Customized pricing quotes

IoT telemetry

### Example architecture



### USE CASE 3

## Pub/Sub or Message Broker

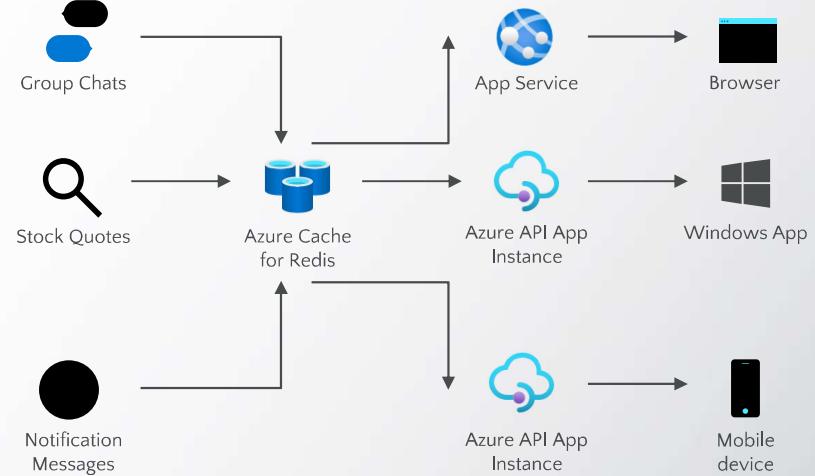
### Use case examples

Publishing news, financial data, or application updates to users

Handling chat messages

Communication between microservices

### Example architecture



#### USE CASE 4

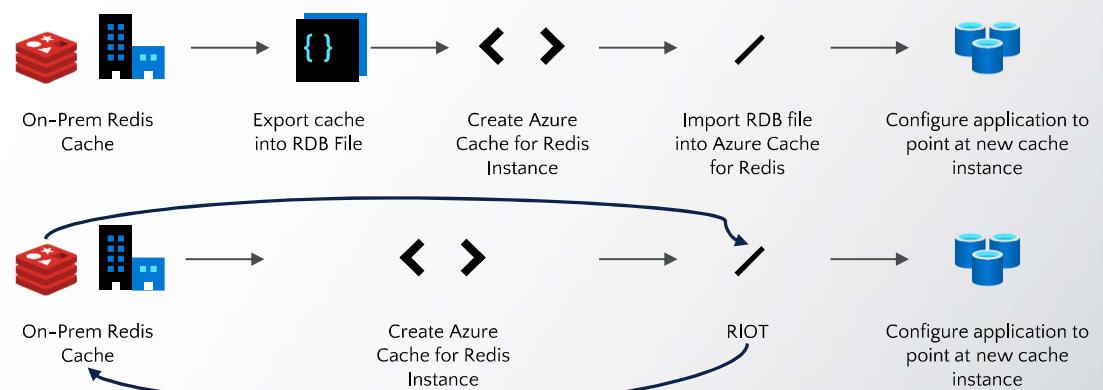
## Migrate Redis to the cloud

### Use case examples

Migrating application from on-premises to the cloud

Modernizing IaaS application through new PaaS services

### Example architecture

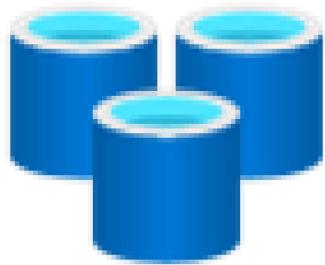


# Azure Cache for Redis Enterprise



# Azure Cache for Redis Offerings

## Basic, Standard, and Premium Tiers



Based on open-source Redis with core features like replication, clustering, and network isolation.

**Best for:** general purpose and dev/test usage

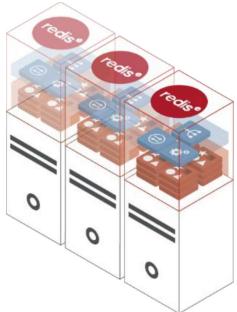
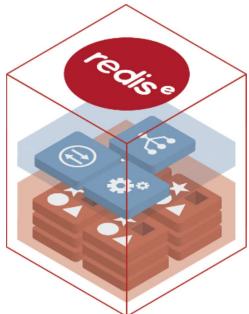
## Enterprise, and Enterprise Flash Tiers



Based on Redis Enterprise with advanced features like active geo-replication, 99.99%+ availability, and Redis Modules

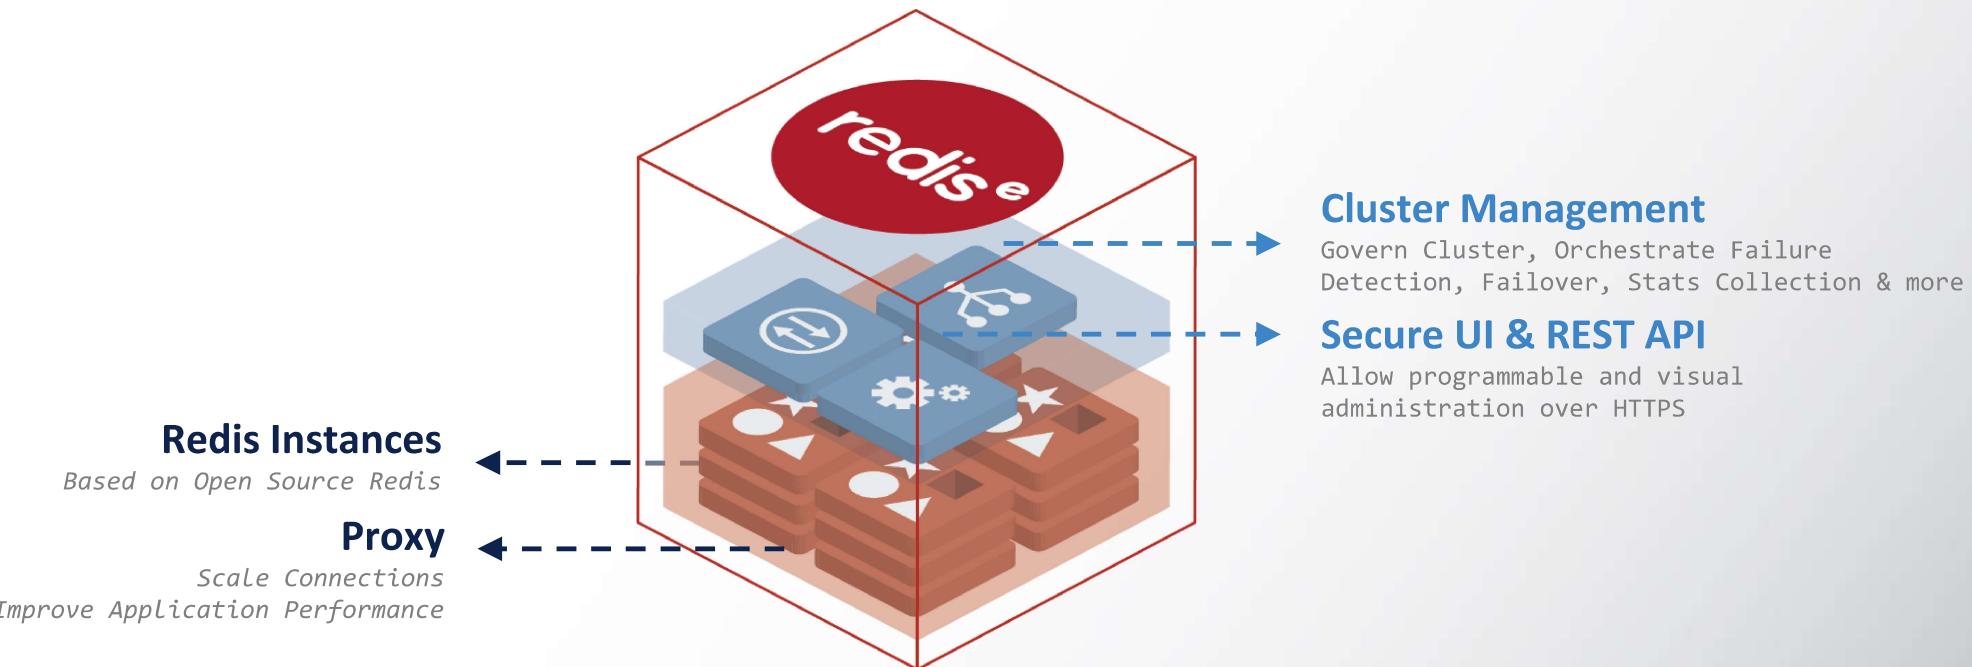
**Best for:** enterprise-scale and advanced usage

# Redis Enterprise Terminology

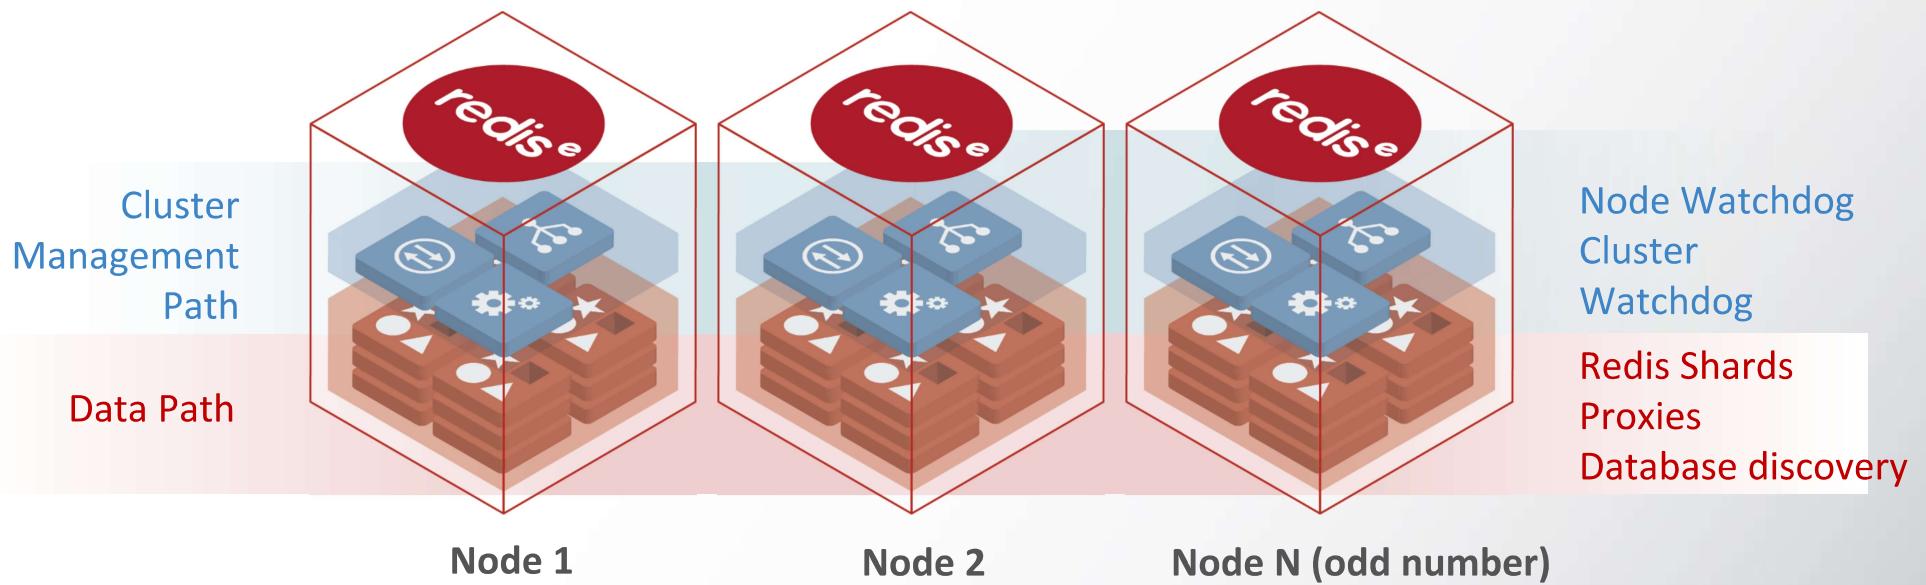


- **Redis Instance** – single-threaded Redis OSS database
  - Commonly referred to as a **shard** with either a role as a master or slave
- **Redis Enterprise Database** – logical entity that manages your entire dataset across multiple Redis Instances
  - not to be confused with the OSS database inside every Redis instance
  - multi-tenant, maximizing infrastructure utilization (reduces overall TCO)
- **Redis Enterprise Nodes** – physical servers, virtual machines, containers, and/or cloud instances
  - on which the Redis Enterprise installation package is installed and configured
  - each node is a container capable of running multiple Redis Instances (i.e., shards)
- **Redis Enterprise Cluster** – set of Redis Enterprise nodes pooling resources
  - allows you to create any number of Redis Enterprise nodes to scale up/out

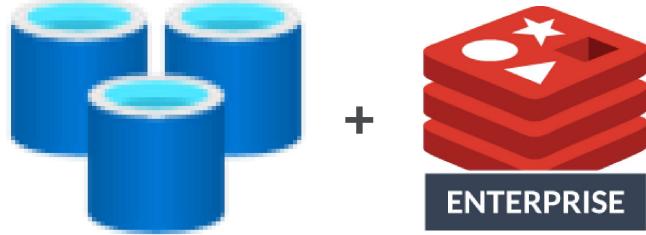
# Redis Enterprise Node



# Shared-Nothing Symmetric Architecture



# Azure Cache for Redis, Enterprise Tiers



Best of both worlds:  
The advanced functionality of Redis Enterprise plus the management and integration of Azure Cache for Redis

## Powerful Features

- Active-Geo Replication
- RediSearch
- RedisBloom
- RedisTimeSeries
- Redis on Flash
- Up to 99.999% availability
- Redis 6.0
- Azure Private Link

## Joint Support

Microsoft provides round-the-clock support and takes the first call.

Direct line to experts at Redis Labs for escalations.

## Streamlined Administration

- Integrated billing
- Azure spend commitments
- Azure security integration
- Azure monitoring tools
- Configuration through the Azure Portal, CLI, ARM, and Terraform

## Redis Enterprise Redis on Flash



Running Redis on high speed NVMe flash storage



Extends your DRAM capacity with SSD and persistent memory



10X larger cache sizes at a lower price per GB



# RedisBloom Module

Extremely efficient probabilistic data structures

## Features

- Includes data structures that trade perfect accuracy for memory efficiency and low latency:
  - **Bloom and Cuckoo filters:** determine if an item is definitely not in a set, or probably is in a set.
  - **Count-min sketch:** determine the frequency of events in a stream.
  - **Top-K:** Find the  $k$  most frequently seen items

## Use-cases

- Checking if usernames are taken
- De-duplication
- Detecting transaction patterns
- Monitoring/alerts
- Leaderboards

## Applications

- Fraud detection
- User management
- Advertising





# RediSearch Module

Real-time secondary index and search engine, built into Redis

## Features

- Extremely fast
- Incremental indexing without performance loss
- Advanced search functionality
  - Prefix
  - Fuzzy
  - Phonetic
  - Stemming
  - Synonyms
  - Autocomplete
  - Geo indexing and filtering

## Use-cases

- Secondary Index
- Full-text search
- Query engine

## Applications

- Enterprise Search
- Real-time inventory
- Customer service & insights
- External database indexing

# RedisTimeSeries Module

High-throughput time-series database



## Features

- Handles extremely high volume of data inserts
- Downsampling/compaction
- Time-based queries (e.g., start time & end time)
- Configurable retention period
- Labels for secondary indexing
- Aggregated Queries
  - Min, Max, Avg, Sum Range, Count, First, Last
  - STD.P, STD.S, Var.P, Var. S

## Use-cases

- Time-series database
- High-throughput data ingestion

## Applications

- IoT telemetry
- Application monitoring
- Anomaly detection

# Active Geo-Replication



Simultaneous read & write operations on multiple globally distributed instances



Worldwide scale with the same local latency



Greater resiliency and data consistency across the globe



# Industry-Leading Availability

**99.9%**

Default

Built-in replication with  
automatic failover to replica  
node

**99.99%**

Multi-AZ

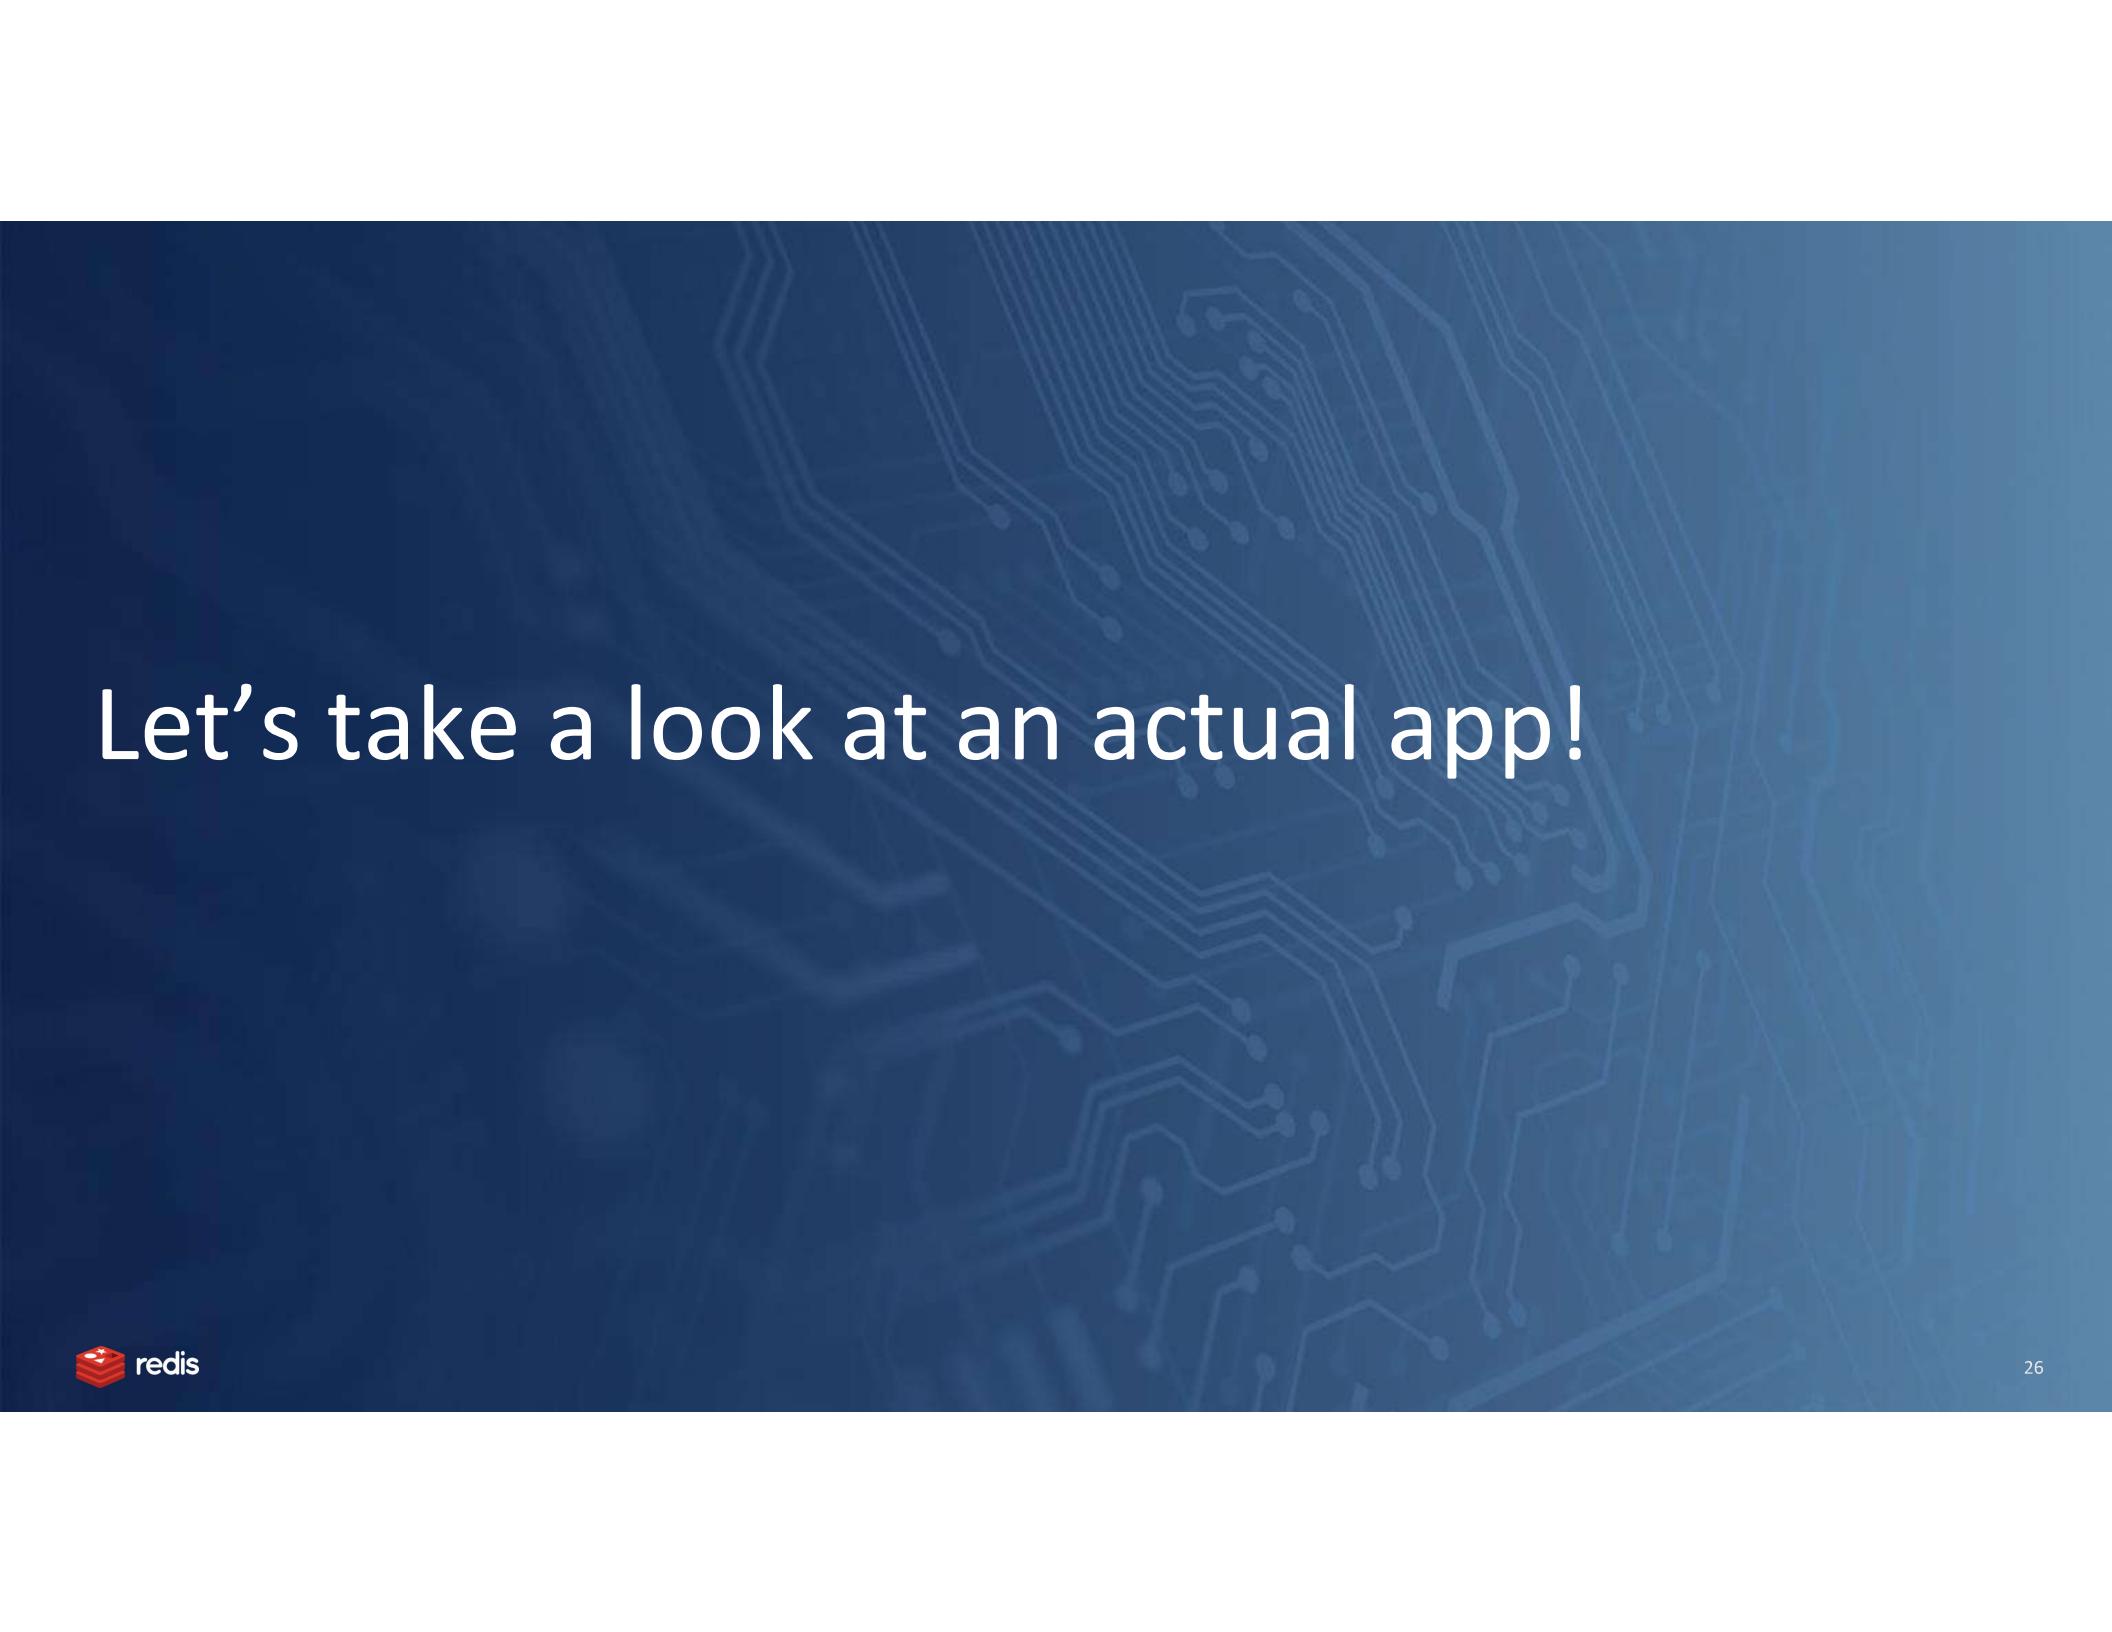
Placing replicas in different  
availability zones

Up to  
**99.999%**

Active Geo-Replication

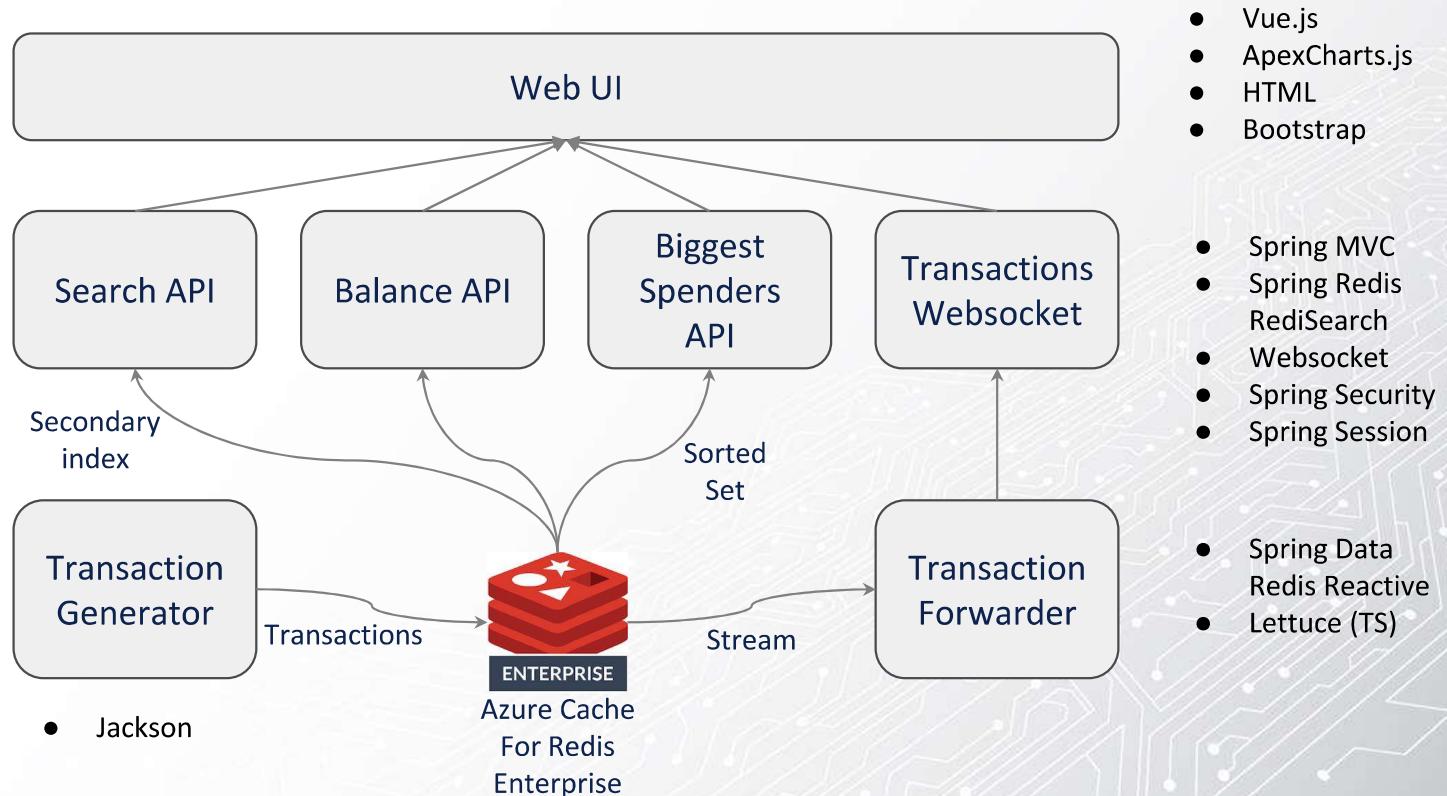
Adding multi-primary writes  
across regions

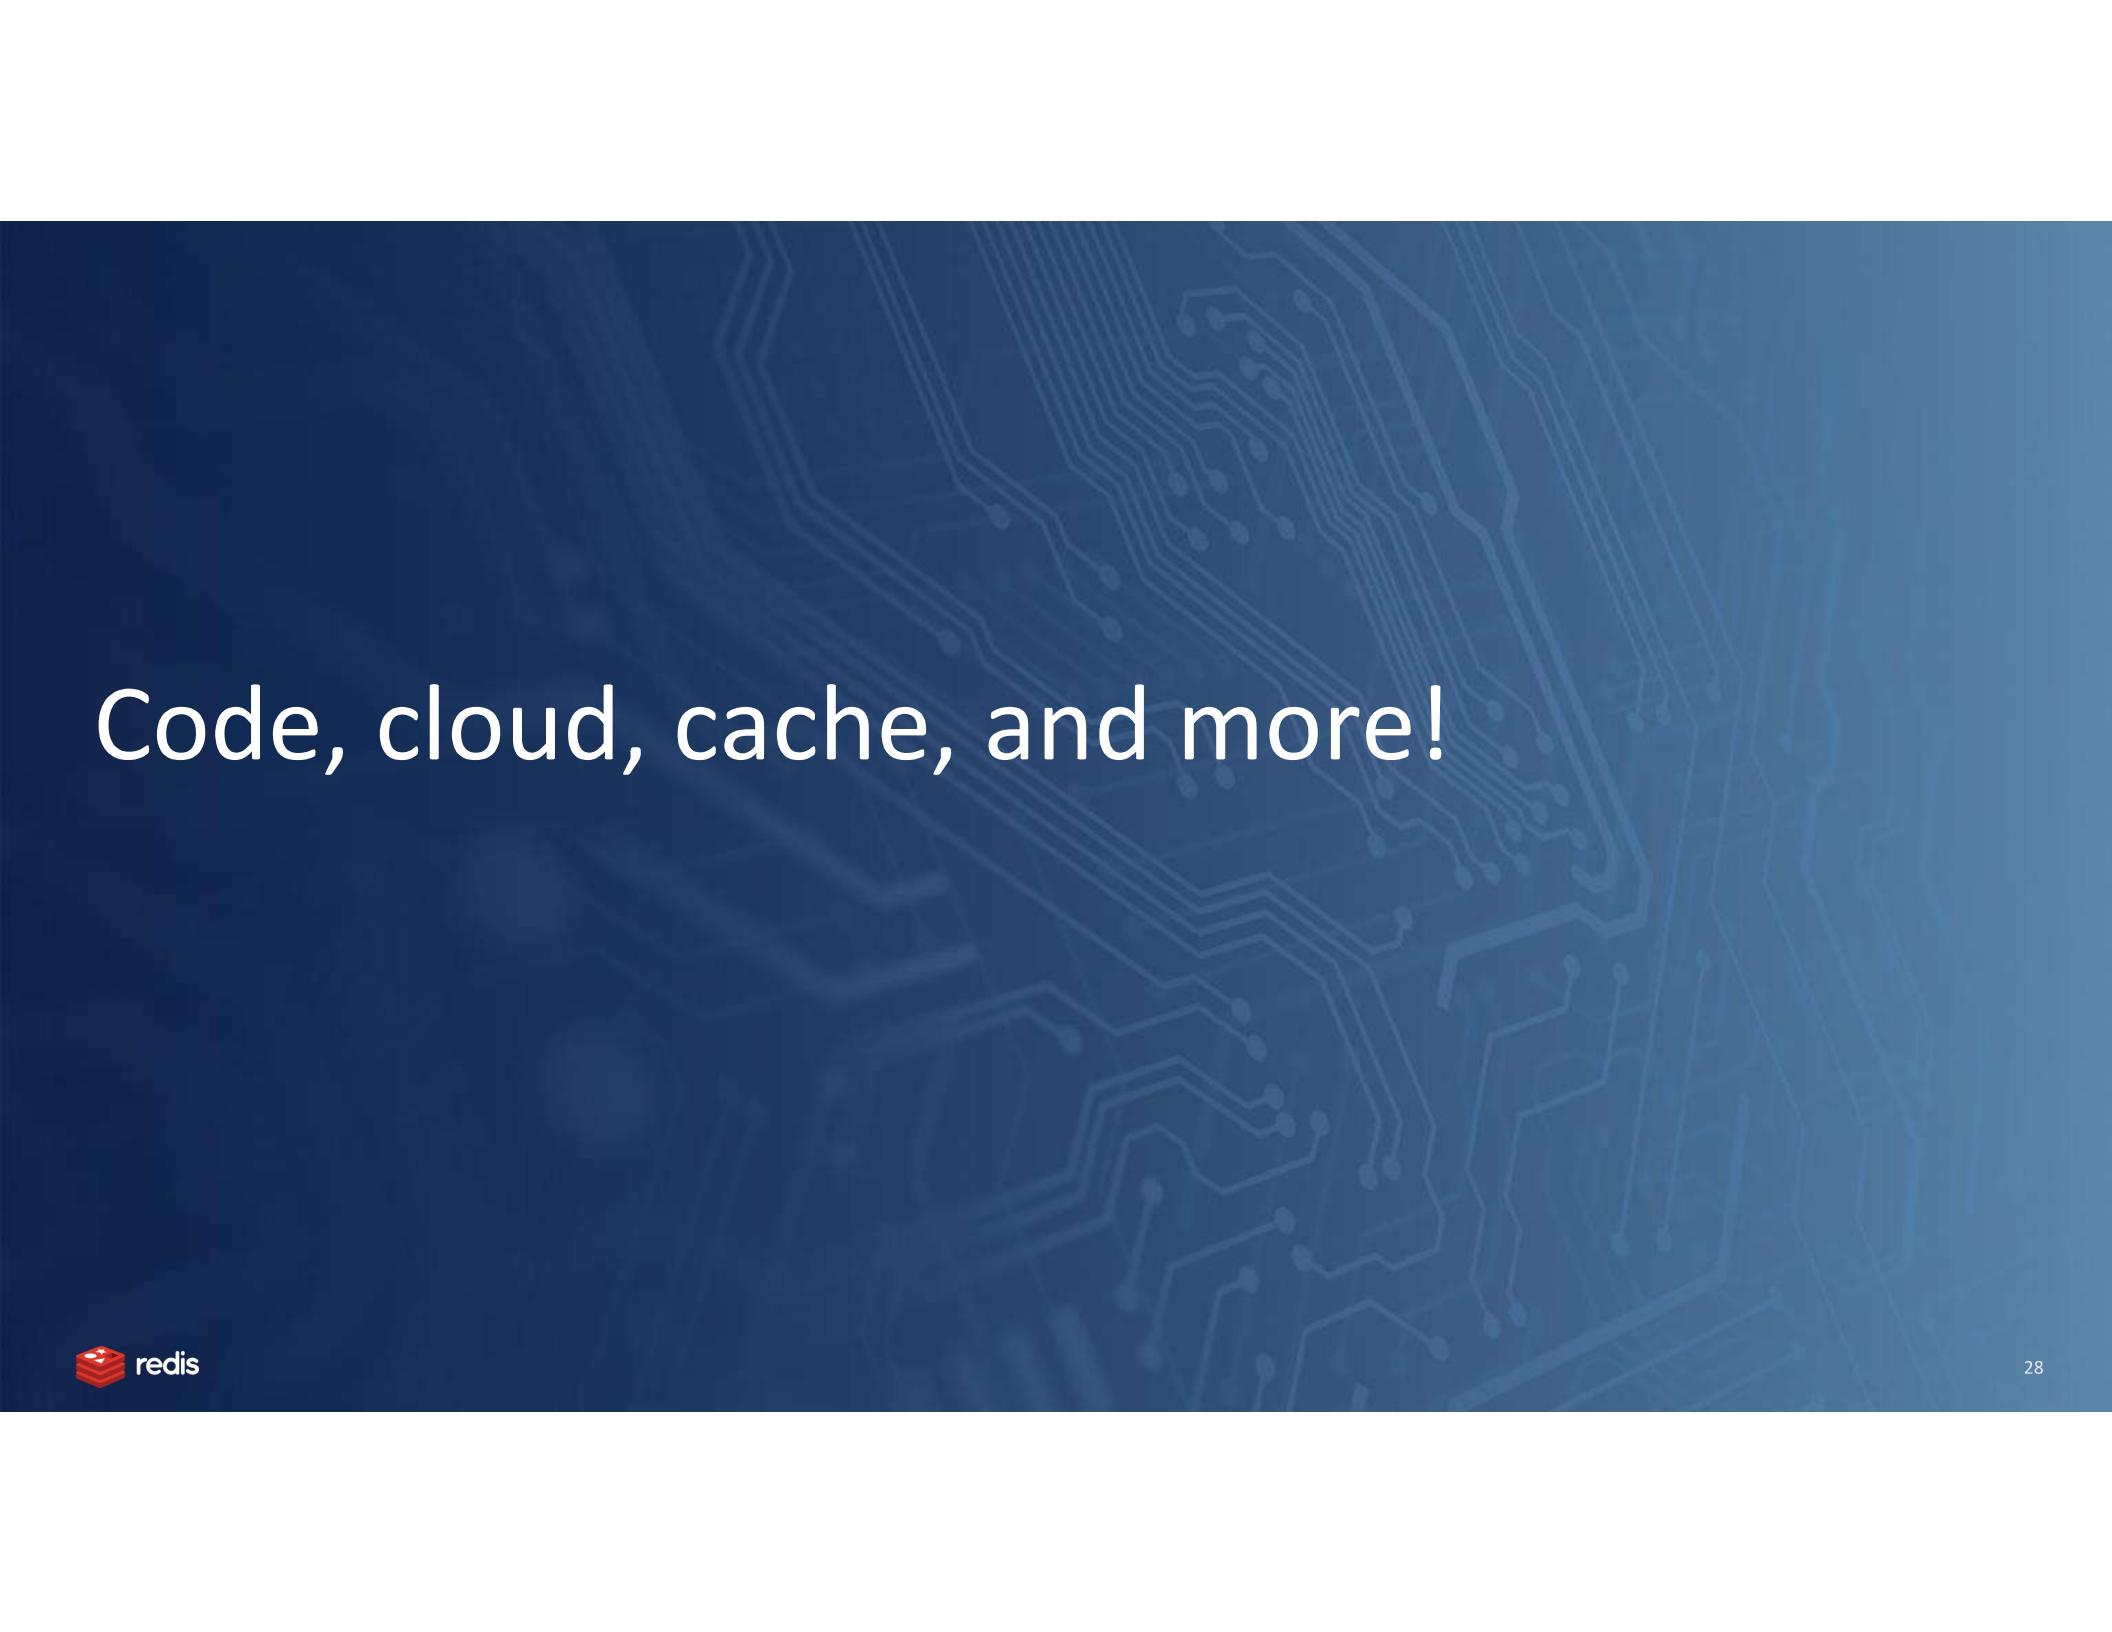


A dark blue background featuring a faint, glowing circuit board pattern with light blue lines and nodes.

# Let's take a look at an actual app!

# Application architecture



A dark blue background featuring a faint, glowing circuit board pattern that curves from the bottom right towards the top left.

# Code, cloud, cache, and more!

# Thank You

